

Exercise 2: Related Processes and Inter-Process Communication via Unnamed Pipes

Operating Systems UE 2019W

Thomas Puchinger

Technische Universität Wien
Computer Engineering
Cyber-Physical Systems

2019-11-28

Content

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of stdin/stdout

Pitfalls

Summary

Related Processes

- ▶ Create a process (fork)
- ▶ Load a new program into a process's memory (exec)
- ▶ Wait on a process's termination (wait)

IPC via Unnamed Pipes

- ▶ (Unnamed) pipe = unidirectional communication channel
- ▶ Communication between **related** processes

Related
Processes

Process
Properties

Interface

Process
Creation

Program
Execution

Process
Termination

Waiting on a
Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of
stdin/stdout

Pitfalls

Summary

Why should we create processes?

- ▶ Divide up a task
 - ▶ Simpler application design
 - ▶ Greater concurrency

Example

A server listens to client requests. The server process starts a new process to handle each request and continues to listen for further connections.

The server can handle several client requests simultaneously.

Process Hierarchy

T. Puchinger

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of stdin/stdout


Pitfalls

Summary

- ▶ Every process has a parent process
- ▶ Exception: init process (init, systemd)
- ▶ Every process has a unique ID (pid_t)
- ▶ Show process hierarchy: `ps tree(1)`

```
systemd-+-ModemManager---2*[{ModemManager}]
        |-NetworkManager--dhclient
        |   '-2*[{NetworkManager}]
        |-abrt-dbus---{abrt-dbus}
        |-2*[abrt-watch-log]
        |-abrt
        |-acpid
        |-agetty
        |-alsactl
        |-atd
        |-auditd-+-audispd-+-sedispatch
        |         |         '-{audispd}
        |         '-{auditd}
        |-automount---7*[{automount}]
        |-avahi-daemon---avahi-daemon
        |-chronyd
        |-colord---2*[{colord}]
        |-crond
        |-cupsd
        |-dbus-daemon
        |-dnsmasq---dnsmasq
        |-firewalld---{firewalld}
        .
        .
```

Memory Layout of a Process



Properties of a Process in Linux

State Running, waiting, ...

Scheduling Priority, CPU time, ...

Identification PID, owner, group, ...

Memory Management Pointer to MMU information

Signals Mask, pending

Process Relations Parents, siblings

Properties of a Process in Linux

Process Control Block Register, PC, status, page table info

Kernel Stack

File description table

Permissions, Accounting Information

Timer Management


Inter-Process communication

See `struct task_struct` in `sched.h`

Interface

fork / exec / exit / wait


- ▶ **fork(2)** – creates a process (copies the process image)
- ▶ **exec(3)** – loads a program (replaces the process image of a process with a new one)
- ▶ **exit(3)** – exits a process
- ▶ **wait(2)** – awaits the exit of child processes



Process Creation

fork

- ▶ Creates a new process
- ▶ New process is an identical copy of the calling process – except PID, pending signals, ...
- ▶ Calling process is the **parent** of the created process, the **child** – processes are **related**
- ▶ Both processes run parallel and execute the same program (from the **fork** call on)



Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes


Redirection of stdin/stdout

Pitfalls


Summary

Process Creation

Before fork()



After fork()



Process Creation

fork

- ▶ Create the process

```
#include <unistd.h>

pid_t fork(void);
```

- ▶ Distinguish between parent and child
via return value of `fork`
 - 1 On error
 - 0 In the child process
 - >0 In the parent process

Process Creation

Example

```
pid_t pid = fork();

switch (pid) {
    case -1:
        fprintf(stderr, "Cannot fork!\n");
        exit(EXIT_FAILURE);

    case 0:
        // child tasks
        ...
        break;
    default:
        // parent tasks
        ...
        break;
}
```

Process Creation

Child

Child inherits from parent:

- ▶ Opened files (common access!)
- ▶ File buffers
- ▶ Signal handling
- ▶ **Current** values of variables

But:

- ▶ Variables are local to process (no influence)
- ▶ Signal handling can be re-configured
- ▶ Communication (IPC) via pipes, sockets, shared memory,
...

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of stdin/stdout


Pitfalls

Summary

Program Execution

exec

- ▶ Load a new program into a process's memory
- ▶ Executes **another** program
- ▶ In the **same** process (PID remains the same)



Program Execution

exec Family¹

```
int execl(const char *path, const char *arg, ...);
int execlp(const char *file, const char *arg, ...);

int execl_e(const char *path, const char *arg, ...,
             char *const envp[]);

int execv(const char *path, char *const argv[]);
int execvp(const char *file, char *const argv[]);

int fexecve(int fd, char *const argv[],
             char *const envp[]);
```

¹Frontend of `execve(2)`

Program Execution

exec Family

- ▶ `exec*p` – searching the environment variable `$PATH` for the program specified
- ▶ `execle` – environment² can be changed
- ▶ `fexecve` – accepts file descriptor (instead of path)
- ▶ `execl*` – variable number of arguments
- ▶ `execv*` – arguments via array

Note Argument Passing!

- ▶ 1st argument is the program's name (`argv[0]`)!
- ▶ Last argument must be a **NULL** pointer!

²FYI: `environ(7)`

Program Execution

Example: `execv()`, `execvp()`

```
#include <unistd.h>
```

```
char *cmd[] = { "ls", "-l", (char *) 0 };
```

```
execv("/bin/ls", cmd);
```

```
execvp("ls", cmd);
```

```
fprintf(stderr, "Cannot exec!\n");
```

```
exit(EXIT_FAILURE);
```

Program Execution

Example: `execl()`, `execlp()`

```
#include <unistd.h>
```

```
execl("/bin/ls", "ls", "-l", NULL);  
// or  
execlp("ls", "ls", "-l", NULL);
```

```
fprintf(stderr, "Cannot exec!\n");  
exit(EXIT_FAILURE);
```

Attention - this is not working:

```
execl("/bin/ls", "ls -l", NULL);
```

```
int a = 1;  
execl("myprog", "myprog", "-a", a, NULL);  
// e.g., use a char-buffer and snprintf(3)
```

Process Termination

exit

T. Puchinger

Related Processes

Process
Properties

Interface

Process
Creation

Program
Execution

Process
Termination

Waiting on a
Child Process

Pitfalls

Debugging

IPC


Pipes

Redirection of
stdin/stdout

Pitfalls

Summary

- ▶ Terminates a process (normally)
- ▶ Termination status can be read by parents
- ▶ Actions performed by `exit()`
 - ▶ Flush and close stdio stream buffers
 - ▶ Close all open files
 - ▶ Delete temporary files (created by `tmpfile(3)`)
 - ▶ Call exit handlers (`atexit(3)`)



Process Termination

exit

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of stdin/stdout

Pitfalls

Summary

- ▶ Terminate a process normally

```
#include <stdlib.h>
```

```
void exit(int status);
```

- ▶ Status: 8 bit (0-255)
- ▶ By convention
 - ▶ `exit(EXIT_SUCCESS)` – process completed successfully
 - ▶ `exit(EXIT_FAILURE)` – error occurred
- ▶ More return values
 - ▶ BSD: `sysexits.h`
 - ▶ <http://tldp.org/LDP/abs/html/exitcodes.html>

Related
Processes

Process
Properties

Interface

Process
Creation

Program
Execution

Process
Termination

Waiting on a
Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of
stdin/stdout


Pitfalls

Summary

Waiting on a Child Process

wait

- ▶ Wait until a child process terminates
- ▶ Returns the PID and status of the terminated child



Waiting on a Child Process

wait

- ▶ Wait for a child to terminate

```
#include <sys/wait.h>

pid_t wait(int *status);
```

- ▶ `wait()` blocks³ until a child terminates or on error
- ▶ Return value
 - ▶ PID of the terminated child
 - ▶ -1 on error (→ `errno`, e.g., `ECHILD`)
- ▶ Status includes exit value and signal information
 - ▶ `WIFEXITED(status)`, `WEXITSTATUS(status)`
 - ▶ `WIFSIGNALED(status)`, `WTERMSIG(status)`
 - ▶ See `wait(2)`

³ ≠ busy waiting

Waiting on a Child Process

Zombies and Orphans

- ▶ UNIX: Terminated processes remain in the process table
- ▶ No more space in process table → no new process can be started!
- ▶ After `wait()` the child process is removed from the process table

Zombie Child terminates, but parent didn't call `wait` yet

- ▶ State of the child is set to "zombie"
- ▶ Child remains in process table until parent calls `wait`

Orphan Parent terminates before child

- ▶ Child gets an **orphan** and is inherited to the `init` process
- ▶ When an orphan terminates, the `init` process removes the entry in the process table

Waiting on a Child Process

Example

```
#include <sys/wait.h>

int status;
pid_t child_pid, pid;
...
while ((pid = wait(&status)) != child_pid)
{
    if (pid != -1) continue;
    // other child
    if (errno == EINTR) continue;
    // interrupted
    fprintf(stderr, "Cannot wait!\n");
    exit(EXIT_FAILURE);
}

if (WEXITSTATUS(status) == EXIT_SUCCESS) {
    ...
}
```


Waiting on a Child Process

waitpid

- ▶ Wait on a **specific** child process

```
#include <sys/wait.h>
```

```
pid_t waitpid(pid_t pid, int *status, int options);
```

- ▶ Examples

```
waitpid(cid, &status, 0);  
    // waits on a child process with PID 'cid'
```

```
waitpid(-1, &status, 0);  
    // equivalent to wait
```

```
waitpid(-1, &status, WNOHANG);  
    // does not block
```

Notification

on Termination of a Child

If parent should not block

- ▶ Synchronous
 - ▶ `waitpid(-1, &status, WNOHANG)`
 - ▶ Returns exit status when a child terminates
 - ▶ Repeating calls → polling
- ▶ Asynchronous
 - ▶ Signal `SIGCHLD` is sent to the parent process whenever one of its child processes terminates
 - ▶ Catch by installing a signal handler (`sigaction`)
 - ▶ Call `wait` in the signal handler

Pitfalls

T. Puchinger

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of stdin/stdout

Pitfalls

Summary

```
int main(int argc, char **argv)
{
    fprintf(stdout, "Hello");

    (void) fork();
    return 0;
}
```

Output: "HelloHello"

Why?

```
int main(int argc, char **argv)
{
    fprintf(stdout, "Hello");
    fflush(stdout);
    (void) fork();
    return 0;
}
```

Output: "Hello"

→ for all opened streams

Debugging

gdb

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of stdin/stdout

Pitfalls

Summary

- ▶ Before fork is executed:
`set follow-fork-mode [child|parent]`

Example

```
$ gdb -tui ./forktest
(gdb) break main
(gdb) set follow-fork-mode child
(gdb) run
(gdb) next
(gdb) :
(gdb) continue
(gdb) quit
```

**Related
Processes**

Process
Properties

Interface

Process
Creation

Program
Execution

Process
Termination

Waiting on a
Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of
stdin/stdout

Pitfalls

Summary

Inter-Process Communication

Recall

So far:

- ▶ Signals (e.g., to synchronise between parent and child)
- ▶ Sockets

New:

- ▶ Pipes

Pipes

Overview


(Unnamed) Pipe

= unidirectional data channel

= enables communication between **related** processes

► Example

```
$ ls | wc -l
```



- Access to read and write end of the pipe via file descriptors
- Pipe is an unidirectional byte stream
- Buffered
- Implicit synchronisation

Pipes

Create

T. Puchinger

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of stdin/stdout

Pitfalls

Summary

- ▶ Create a pipe

```
#include <unistd.h>

int pipe(int pipefd[2]);
```


- ▶ File descriptors of read and write end are returned in specified integer array `pipefd`
 - ▶ `pipefd[0]` – read end
 - ▶ `pipefd[1]` – write end
- ▶ Close unused ends
- ▶ Use read/write end via stream-IO (`fdopen`, etc.)
- ▶ A child process inherits the pipe → common access

Unnamed Pipes

Illustration

`pipe;`

Parent process



Related Processes

Process Properties
Interface
Process Creation
Program Execution
Process Termination
Waiting on a Child Process
Pitfalls
Debugging

IPC

Pipes
Redirection of stdin/stdout
Pitfalls


Summary

Unnamed Pipes



Illustration

`pipe; fork;`

Parent process



Child process



Related Processes

- Process Properties
- Interface
- Process Creation
- Program Execution
- Process Termination
- Waiting on a Child Process
- Pitfalls
- Debugging

IPC

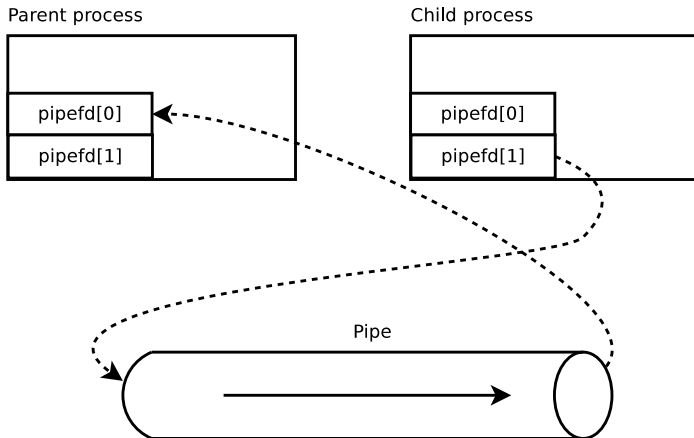
- Pipes
- Redirection of stdin/stdout
- Pitfalls

Summary

Unnamed Pipes

Illustration

`pipe; fork; close unused ends;`



Related Processes

- Process Properties
- Interface
- Process Creation
- Program Execution
- Process Termination
- Waiting on a Child Process
- Pitfalls
- Debugging

IPC

- Pipes
- Redirection of stdin/stdout
- Pitfalls

Summary

Unnamed Pipes

Implicit Synchronisation

- ▶ read blocks on empty pipe
- ▶ write blocks on full pipe
- ▶ read indicates **end-of-file** if all write ends are closed (return value 0)
- ▶ write creates signal **SIGPIPE** if all read ends are closed (if signal ignored/handled: write fails with **errno EPIPE**)

Therefore...

... close unused ends, to get this behaviour (end-of-file and **SIGPIPE/EPIPE**).

Besides, the kernel removes pipes with all ends closed.

Related
Processes

Process
Properties

Interface

Process
Creation

Program
Execution

Process
Termination

Waiting on a
Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of
stdin/stdout

Pitfalls

Summary

Redirection of stdin/stdout

Why?

- ▶ Main application: pipes
- ▶ Example: shell redirection of `stdin` and `stdout`

Scenario:

- ▶ A process may be forked or not
→ uses standard IO
- ▶ A parent process forks and executes another program
- ▶ Parent usually wants to use the child's output
→ redirect `stdin` (file descriptor 0, `STDIN_FILENO`)
and/or `stdout` (file descriptor 1, `STDOUT_FILENO`) in new
process

Redirection of stdin/stdout

File Descriptor Table

T. Puchinger

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

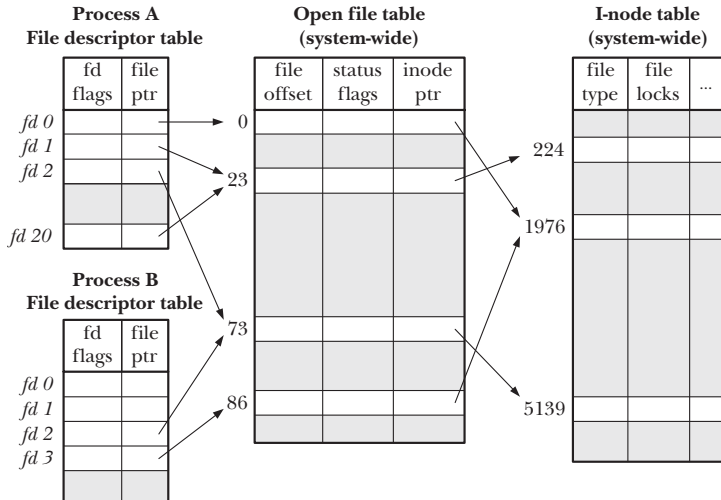
IPC

Pipes

Redirection of stdin/stdout

Pitfalls

Summary



Redirection of stdin/stdout

Approach

T. Puchinger

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of stdin/stdout

Pitfalls

Summary

- ▶ Close file descriptors for standard I/O (stdin, stdout)
- ▶ Duplicate opened file descriptor (e.g., a pipe's end) to the closed one

```
#include <unistd.h>
```

```
int dup(int oldfd);
```

```
int dup2(int oldfd, int newfd);
```

- ▶ Close duplicated file descriptor

Redirection of stdin/stdout

dup / dup2

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of stdin/stdout

Pitfalls

Summary

- ▶ `dup(oldfd)` duplicates file descriptor `oldfd`
 - ▶ New file descriptor uses smallest unused ID
= entry in [file descriptor table](#)
 - ▶ Duplicated file descriptor points to the [same](#) open file description (equal file offset, status flags) → see `open(2)`
- ▶ `dup2(oldfd, newfd)` duplicates `oldfd`
 - ▶ New file descriptor uses ID `newfd`
 - ▶ (Implicitly) closes the file descriptor `newfd` (if necessary)
 - ▶ `newfd` points to the [same](#) open file description like `oldfd`

Redirection of stdin/stdout

Example: redirect stdout to opened file

T. Puchinger

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of stdin/stdout

Pitfalls

Summary

Process A
File descriptor table

	fd	flags	file ptr
(stdin)	fd 0		
(stdout)	fd 1		
(stderr)	fd 2		

(stdin) fd 0 → 0
(stdout) fd 1 → 23
(stderr) fd 2 → 30

Open file table
(system-wide)

file offset	status flags	inode ptr

Redirection of stdin/stdout

Example: redirect stdout to opened file

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

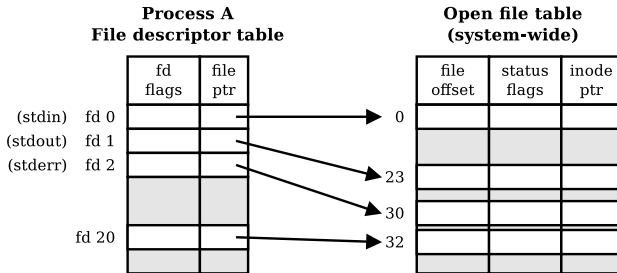
Pipes

Redirection of stdin/stdout

Pitfalls

Summary

open file;



Redirection of stdin/stdout

Example: redirect stdout to opened file

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

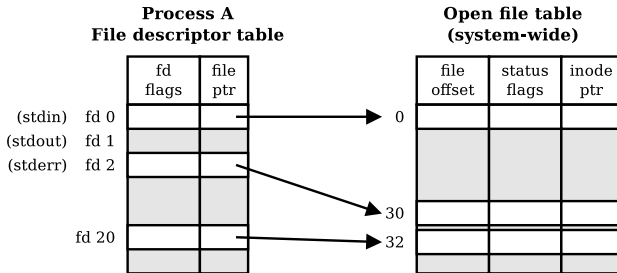
Pipes

Redirection of stdin/stdout

Pitfalls

Summary

open file; close stdout;



Redirection of stdin/stdout

Example: redirect stdout to opened file

T. Puchinger

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

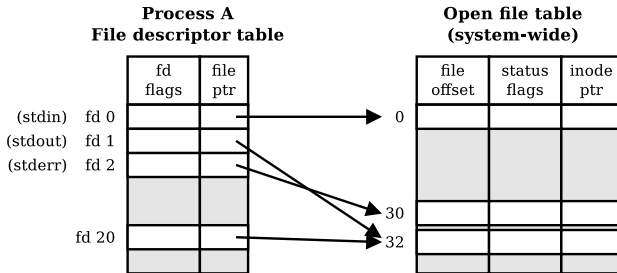
Pipes

Redirection of stdin/stdout

Pitfalls

Summary

open file; close stdout; dup;



Redirection of stdin/stdout

Example: redirect stdout to opened file

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

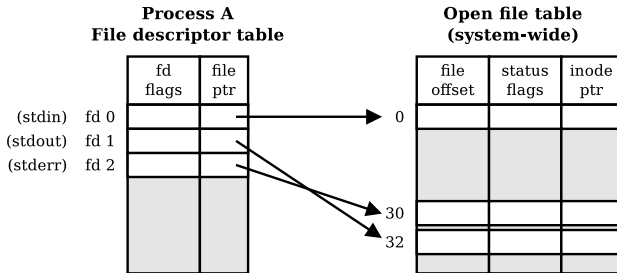
Pipes

Redirection of stdin/stdout

Pitfalls

Summary

open file; close stdout; dup; close file;



Redirection of stdin/stdout

Example: redirect stdout to log.txt

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of stdin/stdout

Pitfalls

Summary

```
#include <fcntl.h>
#include <sys/types.h>
#include <unistd.h>

int fd;

// TODO error handling!

fd = open("log.txt", O_WRONLY | O_CREAT);

dup2(fd,          // old descriptor
      STDOUT_FILENO); // new descriptor

close(fd);

execlp("ls", "ls", NULL);
```

Redirection of stdin/stdout

Example: redirect stdin to pipe

```
// TODO error handling!
```

```
int pipefd[2];  
pipe(pipefd); // create pipe
```

```
pid_t pid = fork();  
switch(pid) {  
:  
case 0: // child counting lines from parent  
    close(pipefd[1]); // close unused write end  
  
    dup2(pipefd[0], // old descriptor - read end  
          STDIN_FILENO); // new descriptor  
  
    close(pipefd[0]);  
  
    execlp("wc", "wc", "-l", NULL);  
    // should not reach this line  
:  
}
```

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of
stdin/stdout

Pitfalls

Summary

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of stdin/stdout

Pitfalls

Summary

- ▶ Pipes are **unidirectional**
- ▶ Bidirectional: two pipes, but ...
 - ▶ Erroneous synchronisation (deadlock, e.g., both processes read from empty pipe)
- ▶ Synchronisation & Buffer
 - ▶ Use `fflush()`
 - ▶ Configure buffer (`setbuf(3)`, `setvbuf(3)`)

Related
Processes

Process
Properties

Interface

Process
Creation

Program
Execution

Process
Termination

Waiting on a
Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of
stdin/stdout

Pitfalls

Summary

Tips for the Exercise

- ▶ Try to parallel the functionality of your program (as much as possible)

Example

DO NOT: The parent first reads all input from a file to an array. It then sends the data within one burst to the child. The child processes the data and outputs the result.

INSTEAD DO: The parent reads line-by-line from a file. Each line is sent to the client immediately. Reading and processing of the lines happens in parallel.

Tips for the Exercise

T. Puchinger

Related Processes

Process Properties

Interface

Process Creation

Program Execution

Process Termination

Waiting on a Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of stdin/stdout

Pitfalls

Summary

- Communicate over pipes (do not exploit inherited memory areas)

Example

DO NOT: The parent reads a file and saves its content into an array and forks a child. The child processes the data from the array.

INSTEAD DO: The parent communicates the data from the file over a pipe.

- However, you may pass options/flags/settings to the child (process). For example, use inherited variable `argv` to set arguments when using `exec`.

Summary

T. Puchinger

Related Processes

Process
Properties

Interface

Process
Creation

Program
Execution

Process
Termination

Waiting on a
Child Process

Pitfalls

Debugging

IPC

Pipes

Redirection of
stdin/stdout

Pitfalls

Summary

- ▶ fork/exec/wait
 - ▶ Start further programs
- ▶ Unnamed Pipes
 - ▶ Communication between related processes
 - ▶ Redirection of stdin/stdout

Material

T. Puchinger

Related Processes

Process Properties
Interface
Process Creation
Program Execution
Process Termination
Waiting on a Child Process
Pitfalls
Debugging

IPC

Pipes
Redirection of stdin/stdout
Pitfalls

Summary

- ▶ Michael Kerrisk: A Linux and UNIX System Programming Handbook, No Starch Press, 2010.
- ▶ man pages: fork(2), exec(3), execve(2), exit(3), wait(3), pipe(2), dup(2)
- ▶ gdb - Debugging Forks:
<https://sourceware.org/gdb/onlinedocs/gdb/Forks.html>